

# Machine Learning HW1

January 19, 2018

## 1 1. Preparation

```
In [1]: import numpy as np
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
plt.style.use('ggplot')

%matplotlib inline

from numpy.linalg import inv
from numpy.random import multivariate_normal, exponential
import statsmodels.formula.api as smf
from scipy import stats

from sklearn.base import RegressorMixin, BaseEstimator, TransformerMixin
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error

from progressbar import ProgressBar
from lazy import lazy
from tabulate import tabulate
```

This AwesomeLinearModel is a class I wrote years ago for Econometrics class. It calculates all the quantities we need for the following questions, so I decided to use it.

```
In [2]: class AwesomeLinearModel(RegressorMixin, BaseEstimator, TransformerMixin):

    def __init__(self):
        super(AwesomeLinearModel, self).__init__()
        self.x_names, self.y_name = None, None
        self.n, self.p = None, None
        self.formula = 'LinearModel( )\n'
        self.fitted = False

    # algebraic properties
```

```

self.X, self.y = None, None
self.XtX = None
self.XtX_inv = None
self.b_hat = None
self.HX = None
self.residuals = None # n*1
self.RSS = None

# statistical properties
self.sigma_sq_mle = None
self.sigma_sq_unbiased = None
self.rse = None
self.cov_b = None
self.var_b = None
self.std_b = None
self.t_stats = None
self.p_val = None

def __repr__(self):
    _repr, rows = self.formula, []
    if not self.fitted:
        return _repr
    for i, s in enumerate(self.x_names):
        rows.append([
            s, self.b_hat[i], self.std_b[i],
            self.t_stats[i], self.p_val[i]])
    _repr += tabulate(
        rows, headers=['', 'Estimate', 'Std.Error',
                       't.Stat', 'p.value'])
    return _repr

def fit(self, X, y):
    """
    Fit the model.
    :param X: features set.
    :param y: response variable set.
    :return:
    """
    n, p = X.shape
    if self.x_names is None:
        self.x_names = ['X%d'%i for i in range(1,p+1)]
        self.y_name = 'Y'
        formula = 'Y ~ %s' % ' + '.join(self.x_names)
        self.formula = 'LinearModel( {} )\n'.format(formula)

# algebraic properties
self.X, self.y = X, y
self.XtX = X.T.dot(X)

```

```

self.XtX_inv = inv(self.XtX)
self.b_hat = self.XtX_inv.dot(X.T.dot(y))
self.HX = X.dot(self.XtX_inv).dot(X.T)
self.residuals = y - X.dot(self.b_hat) # n*1
self.RSS = sum([r[0] * r[0] for r in self.residuals])

# statistical properties
self.sigma_sq_mle = self.RSS / n
self.sigma_sq_unbiased = self.RSS / (n - p)
self.rse = np.sqrt(self.sigma_sq_unbiased)
self.cov_b = self.XtX_inv * self.sigma_sq_unbiased
self.var_b = np.diag(self.cov_b).reshape(p, 1)
self.std_b = np.sqrt(self.var_b).reshape(p, 1)
self.t_stats = self.b_hat / self.std_b
self.p_val = (stats.t.sf(
    np.abs(self.t_stats), n - p) * 2).reshape(p, 1)

self.fitted = True
return self

def predict(self, X_test):
    return X_test.dot(self.b_hat)

@property
def coef(self):
    return self.b_hat

```

## 2 2. Linear Regression with Correlation

```

In [3]: def lr_correlated_variables(rho_range, true_beta,
                                     n_sample=1000, test_portion=0.5):
    # initialization
    lr_model = AwesomeLinearModel()
    mse_path, var_path = [], []

    for rho in rho_range:
        # simulate independent X, epsilon
        X = multivariate_normal(
            np.array([0,0]),
            np.array([[1,rho],[rho,1]]),
            size=n_sample
        )
        epsilon = multivariate_normal(
            np.array([0]),np.array([[1]]),n_sample)
        # generate y
        y = X.dot(beta_dgp) + epsilon

```

```

# train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=test_portion, random_state=42)

# fit model, store the quantities we want
lr_model.fit(X_train, y_train)
y_pred = lr_model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
var_beta1 = lr_model.var_b[0][0]
var_path.append(var_beta1)
mse_path.append(mse)
return rho_range, var_path, mse_path

```

```

In [55]: np.random.seed(42)
# run simulation
rho_range, var_path, mse_path = lr_correlated_variables(
    np.linspace(0,0.9,10),
    np.array([1,1]).reshape(2,1),
    1000, 0.5)

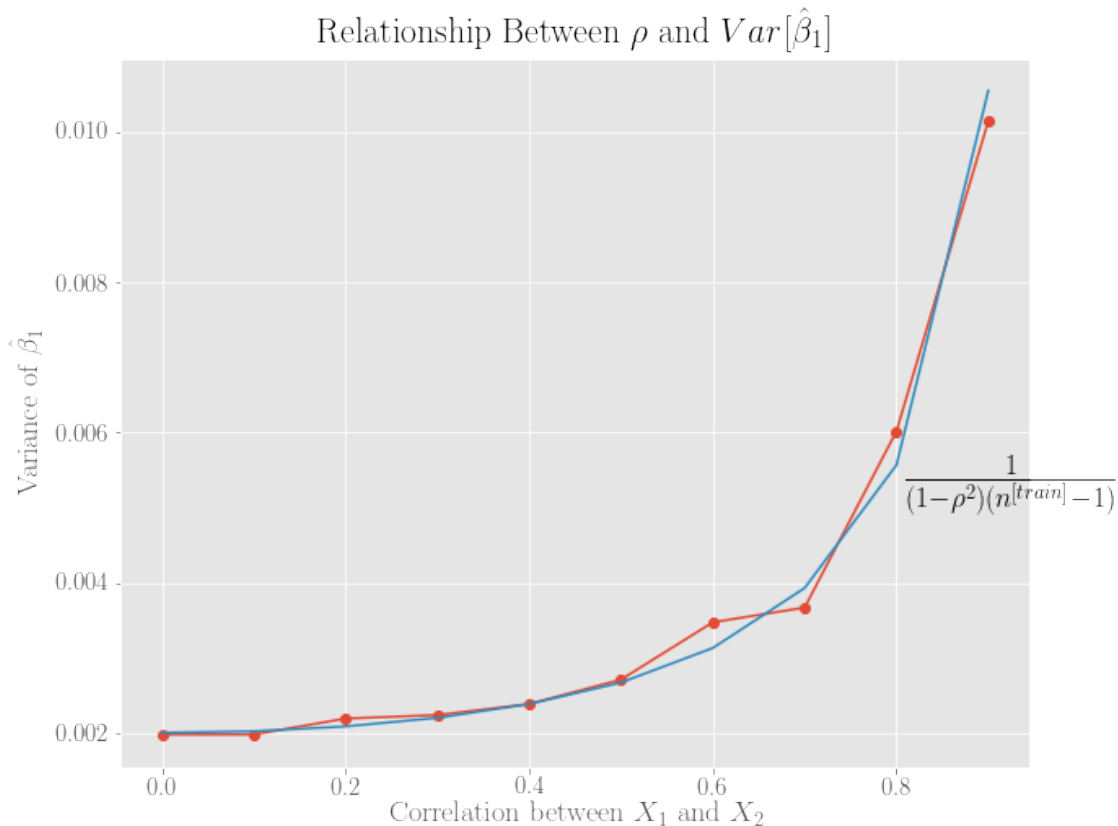
# plot
fig, ax = plt.subplots(1, 1, figsize=(10,8))
ax.plot(rho_range, var_path, 'o-')
ax.plot(rho_range, 1/((1-rho_range**2)*499))
ax.set_xlabel('Correlation between $X_1$ and $X_2$')
ax.set_ylabel('Variance of $\hat{\beta}_1$')
ax.set_title('Relationship Between $\rho$ and $Var[\hat{\beta}_1]$')
plt.rc('text', usetex=True)
plt.rc('font', family='serif', size=15)
plt.text(0.81, 0.005, r'$\frac{1}{(1-\rho^2)(n^{[train]}-1)}$', size=25)

```

```

Out[55]: <matplotlib.text.Text at 0x12057be80>

```



It can be shown that the estimated variance of  $\hat{\beta}_j$  can be equivalently expressed as

$$\hat{Var}[\hat{\beta}_j] = \frac{\hat{\sigma}_\epsilon^2}{(n-1)\hat{Var}[X_j]} \frac{1}{(1-R_j^2)}$$

Which turns out to be

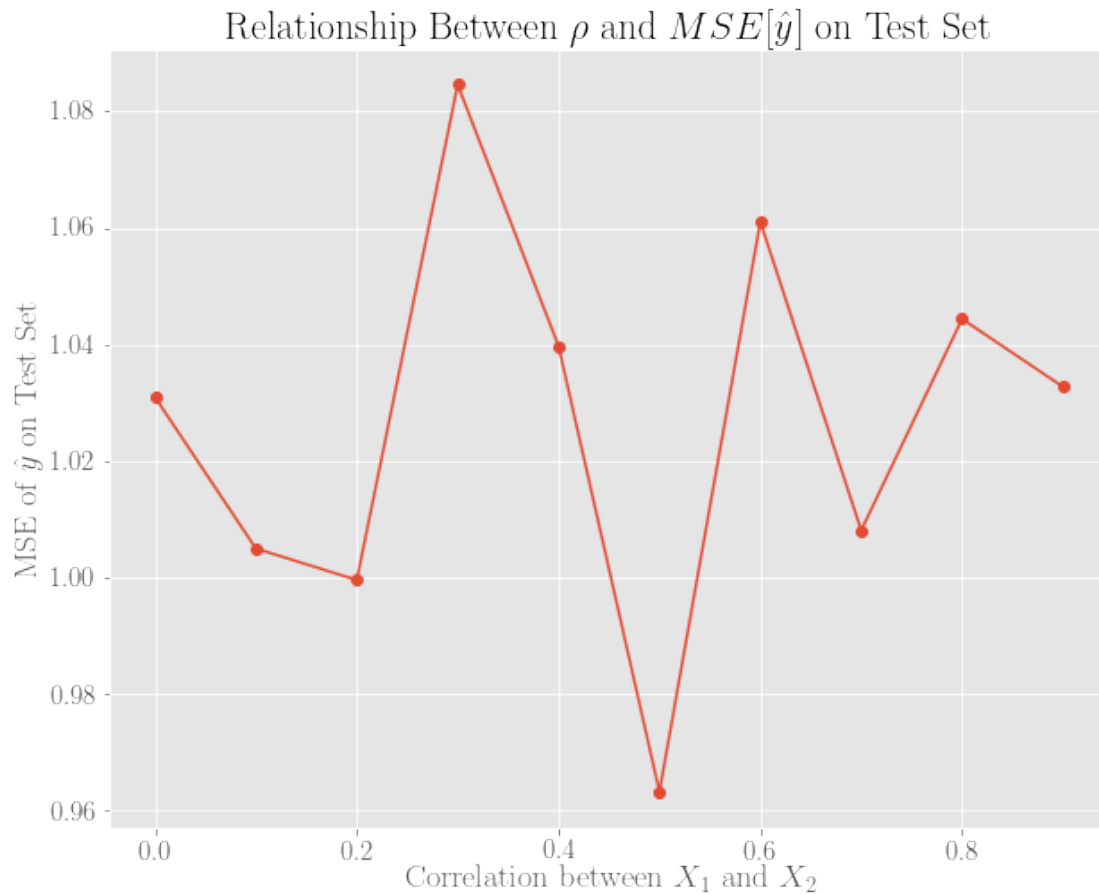
$$\hat{Var}[\hat{\beta}_j] = \frac{1}{(1-\rho^2)(n^{[train]}-1)}$$

in this problem. So we see the simulation results agree to the blue line.

In [51]: *# another plot*

```
fig, ax = plt.subplots(1, 1, figsize=(10,8))
ax.plot(rho_range, mse_path, 'o-')
ax.set_xlabel('Correlation between $X_1$ and $X_2$')
ax.set_ylabel('MSE of $\hat{y}$ on Test Set')
ax.set_title('Relationship Between $\rho$ and $MSE[\hat{y}]$ on Test Set')
```

Out[51]: <matplotlib.text.Text at 0x1200b9320>



There is no strong relationship between  $\rho$  and the prediction error, measured by the mean squared error of  $\hat{y}$  on the test set.

## 2.1 3. Linear Regression in Dimensions

```
In [106]: X = multivariate_normal(
            np.zeros(3),
            np.identity(3),
            10
        )
X[:,0:1]
```

```
Out[106]: array([[ 0.28018636],
                  [ 0.10307429],
                  [-0.779272  ],
                  [-1.23021135],
                  [ 1.21176011],
                  [ 0.16747375],
                  [-0.67725433],
                  [ 1.03296001],
```

```

        [-0.00371641],
        [-0.89865235]])

```

```

In [108]: def lr_high_dimensions(dim_range, beta_dgp, n_epochs=100,
                                n_sample=1000, test_portion=0.5):
    mse_table_test = np.zeros((n_epochs, len(dim_range)))
    mse_table_train = np.zeros((n_epochs, len(dim_range)))
    bar = ProgressBar()
    m = dim_range[-1]
    for j, p in bar(list(enumerate(dim_range))):
        for i in range(n_epochs):
            # simulate independent X, epsilon
            X = multivariate_normal(
                np.zeros(m),
                np.identity(m),
                size=n_sample
            )
            epsilon = multivariate_normal(
                np.array([0]), np.array([[1]]), n_sample)

            # generate y
            y = X.dot(beta_dgp) + epsilon

            # fit model, store the quantities we want
            X_train, X_test, y_train, y_test = train_test_split(
                X[:,0:p], y, test_size=test_portion, random_state=42)
            lr_model = AwesomeLinearModel()
            lr_model.fit(X_train, y_train)
            y_pred_test = lr_model.predict(X_test)
            y_pred_train = lr_model.predict(X_train)
            mse_table_test[i][j] = mean_squared_error(
                y_test, y_pred_test)
            mse_table_train[i][j] = mean_squared_error(
                y_train, y_pred_train)
        mse_path_test = mse_table_test.mean(0)
        mse_path_train = mse_table_train.mean(0)
    return dim_range, mse_path_test, mse_path_train

```

```

In [112]: np.random.seed(42)
           # run simulation
           dim_range, mse_path_test, mse_path_train = lr_high_dimensions(
               range(1,81),
               np.array([4]+79*[0]).reshape(80,1),
               n_epochs=100,
               n_sample=1100,
               test_portion=10/11)

```

```


# plot
fig, ax = plt.subplots(1, 1, figsize=(10,8))
ax.plot(dim_range, mse_path_test, 'o-')
ax.plot(dim_range, mse_path_train, 'o-')
ax.set_xlabel('Model Complexity: \#Variables Included')
ax.set_ylabel('MSE of  $\hat{y}$ ')
ax.set_title('Relationship Between \#Variables and  $MSE[\hat{y}]$  on Training and Test')
ax.legend(('MSE on Test Set', 'MSE on Training Set'))

```

100% (80 of 80) |#####| Elapsed Time: 0:01:00 Time: 0:01:00

Out[112]: <matplotlib.legend.Legend at 0x123b8d710>

Relationship Between #Variables and  $MSE[\hat{y}]$  on Training and Test Set



## 2.2 4. Exponential Prediction

```

In [33]: np.random.seed(42)
n_sample = 1000
y = exponential(1, n_sample)
y_pred_mean = np.ones(n_sample)

```



```
y_pred_med = np.ones(n_sample) * np.log(2)
mean_squared_error(y, y_pred_mean), mean_squared_error(y, y_pred_med)
```

```
Out[33]: (0.94558562609916719, 1.022871024128073)
```

```
In [34]: mean_absolute_error(y, y_pred_mean), mean_absolute_error(y, y_pred_med)
```

```
Out[34]: (0.73302558820985386, 0.68460817591549794)
```

The mean squared error prefers the mean estimator, while the mean absolute error metric prefers the median estimator

```
In [ ]:
```