

HW4

Ze YANG

10/1/2018

R Lab from Ruppert

Problem 1

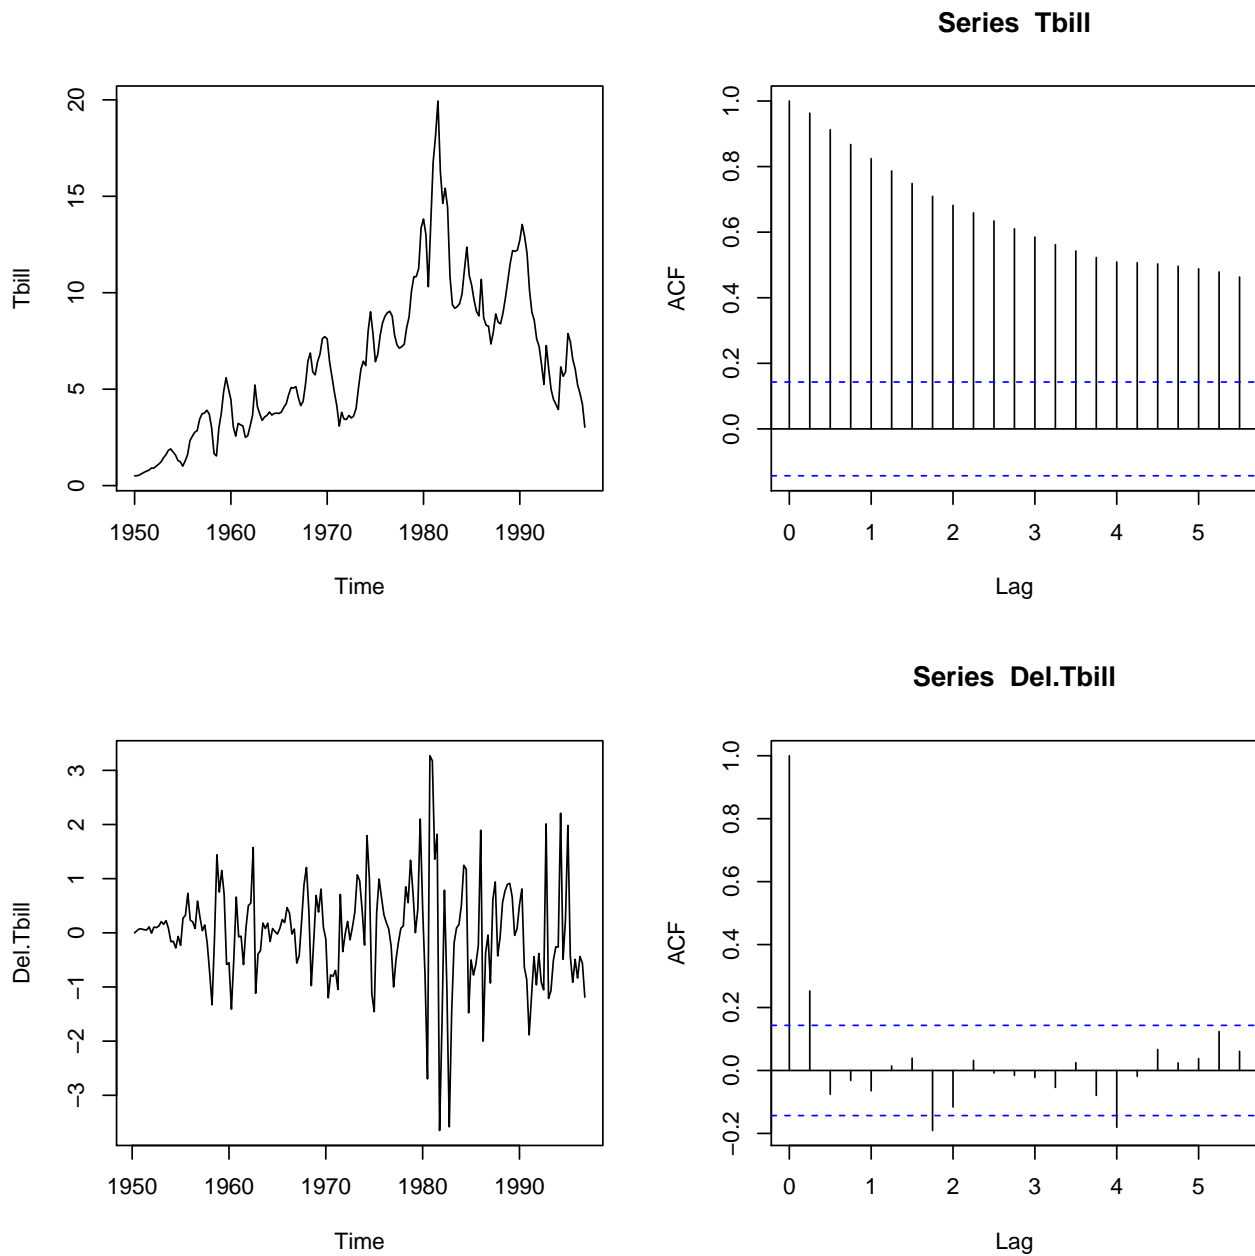
```
data(Tbrate, package="Ecdat")
library(forecast)

## Warning in as.POSIXlt.POSIXct(Sys.time()): unknown timezone 'default/
## America/New_York'

library(tseries)
library(fGarch)

## Loading required package: timeDate
## Loading required package: timeSeries
## Loading required package: fBasics
# r = the 91-day treasury bill rate
# y = the log of real GDP
# pi = the inflation rate
Tbill = Tbrate[,1]
Del.Tbill = diff(Tbill)

par(mfrow=c(2,2))
plot(Tbill)
acf(Tbill)
plot(Del.Tbill)
acf(Del.Tbill)
```



```
adf.test(Tbill)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: Tbill
## Dickey-Fuller = -1.925, Lag order = 5, p-value = 0.6075
## alternative hypothesis: stationary
```

```
adf.test(Del.Tbill)
```

```
## Warning in adf.test(Del.Tbill): p-value smaller than printed p-value
##
## Augmented Dickey-Fuller Test
##
```

```
## data: Del.Tbill
## Dickey-Fuller = -5.2979, Lag order = 5, p-value = 0.01
## alternative hypothesis: stationary
```

Comments - Tbill's plot has a clear trend in it, and the acf plot decays slowly. The ADF test failed to reject null H_0 : the series is nonstationary. Hence we believe it's nonstationary.

- Del.Tbill's plot displays mean-reverting behavior; the acf plot decays fast. The ADF test also reject the nonstationarity hypothesis with confidence level 0.01. So we think it's stationary.
- Del.Tbill exhibits two types of heteroscedasticity patterns: First, it has a volatility clustering behavior. Second, the volatility is correlated to the value of Tbill: the vol's higher in early 1980s when the local mean of Tbill itself is high.

Problem 2

```
garch.model.Tbill = garchFit(formula= ~arma(1,0) + garch(1,0),Del.Tbill)
```

```
##
## Series Initialization:
## ARMA Model:          arma
## Formula Mean:        ~ arma(1, 0)
## GARCH Model:         garch
## Formula Variance:    ~ garch(1, 0)
## ARMA Order:          1 0
## Max ARMA Order:      1
## GARCH Order:         1 0
## Max GARCH Order:     1
## Maximum Order:       1
## Conditional Dist:    norm
## h.start:             2
## llh.start:           1
## Length of Series:    187
## Recursion Init:      mci
## Series Scale:        0.9422364
##
## Parameter Initialization:
## Initial Parameters:   $params
## Limits of Transformations: $U, $V
## Which Parameters are Fixed? $includes
## Parameter Matrix:
##
##      U          V      params includes
## mu    -0.14290725  0.1429072 0.01197006    TRUE
## ar1   -0.99999999  1.0000000 0.25347489    TRUE
## omega 0.00000100 100.0000000 0.10000000    TRUE
## alpha1 0.00000001  1.0000000 0.10000000    TRUE
## gamma1 -0.99999999  1.0000000 0.10000000    FALSE
## delta  0.00000000  2.0000000 2.00000000    FALSE
## skew   0.10000000 10.0000000 1.00000000    FALSE
## shape  1.00000000 10.0000000 4.00000000    FALSE
## Index List of Parameters to be Optimized:
##      mu    ar1  omega alpha1
##      1     2     3     4
## Persistence:          0.1
```

```

##
##
## --- START OF TRACE ---
## Selected Algorithm: nlminb
##
## R coded nlminb Solver:
##
## 0:      456.21058: 0.0119701 0.253475 0.100000 0.100000
## 1:      251.86762: 0.0119754 0.236936 1.02079 0.489709
## 2:      250.90930: 0.0120591 0.473682 0.845081 0.893495
## 3:      246.37888: 0.0132650 0.303699 0.757428 1.00000
## 4:      244.30653: 0.0172433 0.228377 0.691915 1.00000
## 5:      240.23776: 0.0270411 0.161803 0.543304 1.00000
## 6:      236.37169: 0.0367592 0.175592 0.370918 1.00000
## 7:      236.31833: 0.0446835 0.210080 0.425054 0.537516
## 8:      235.93894: 0.0447059 0.288586 0.470834 0.622965
## 9:      235.66475: 0.0444909 0.288456 0.377870 0.704783
## 10:     235.36527: 0.0456870 0.217083 0.405800 0.756581
## 11:     235.19182: 0.0473954 0.248964 0.405265 0.772573
## 12:     235.13214: 0.0513735 0.255349 0.382404 0.810199
## 13:     235.08391: 0.0553971 0.257352 0.383930 0.823996
## 14:     234.96674: 0.0712269 0.253195 0.382822 0.857205
## 15:     234.92037: 0.0846539 0.246053 0.380919 0.862586
## 16:     234.91025: 0.0885175 0.242328 0.380676 0.848471
## 17:     234.90819: 0.0890151 0.241257 0.380809 0.836077
## 18:     234.90814: 0.0886908 0.241545 0.380857 0.834926
## 19:     234.90814: 0.0886149 0.241631 0.380886 0.834829
## 20:     234.90814: 0.0886153 0.241635 0.380889 0.834828
##
## Final Estimate of the Negative LLH:
## LLH: 223.7818      norm LLH: 1.196694
##      mu      ar1      omega      alpha1
## 0.08349652 0.24163450 0.33815662 0.83482811
##
## R-optimhess Difference Approximated Hessian Matrix:
##      mu      ar1      omega      alpha1
## mu      -379.01802 -68.15986 -37.51437 4.76325
## ar1      -68.15986 -209.40200 51.00153 8.23281
## omega    -37.51437 51.00153 -399.68185 -54.31253
## alpha1    4.76325 8.23281 -54.31253 -24.58337
## attr("time")
## Time difference of 0.01015687 secs
##
## --- END OF TRACE ---
##
##
## Time to Estimate Parameters:
## Time difference of 0.05084801 secs
summary(garch.model.Tbill)

##
## Title:
## GARCH Modelling
##

```

```
## Call:
## garchFit(formula = ~arma(1, 0) + garch(1, 0), data = Del.Tbill)
##
## Mean and Variance Equation:
## data ~ arma(1, 0) + garch(1, 0)
## <environment: 0x7fadd0b80830>
## [data = Del.Tbill]
##
## Conditional Distribution:
## norm
##
## Coefficient(s):
##      mu      ar1      omega      alpha1
## 0.083497 0.241635 0.338157 0.834828
##
## Std. Errors:
## based on Hessian
##
## Error Analysis:
##      Estimate Std. Error t value Pr(>|t|)
## mu      0.08350    0.05391   1.549 0.121395
## ar1     0.24163    0.07280   3.319 0.000902 ***
## omega   0.33816    0.06145   5.503 3.73e-08 ***
## alpha1  0.83483    0.24295   3.436 0.000590 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log Likelihood:
## -223.7818      normalized: -1.196694
##
## Description:
## Mon Oct  1 15:38:53 2018 by user:
##
##
## Standardised Residuals Tests:
##
##                               Statistic p-Value
## Jarque-Bera Test      R      Chi^2 26.96616 1.394352e-06
## Shapiro-Wilk Test     R      W      0.957289 1.957134e-05
## Ljung-Box Test        R      Q(10) 10.8275 0.3711143
## Ljung-Box Test        R      Q(15) 13.10815 0.5939444
## Ljung-Box Test        R      Q(20) 16.11144 0.7096868
## Ljung-Box Test        R^2 Q(10) 12.18313 0.2729873
## Ljung-Box Test        R^2 Q(15) 14.31743 0.5016035
## Ljung-Box Test        R^2 Q(20) 17.28754 0.634231
## LM Arch Test          R      TR^2  9.685432 0.6435358
##
## Information Criterion Statistics:
##      AIC      BIC      SIC      HQIC
## 2.436169 2.505284 2.435279 2.464174
garch.model.Tbill@fit$matcoef
##      Estimate Std. Error t value Pr(>|t|)
## mu      0.08349652 0.05390549 1.548943 1.213955e-01
## ar1     0.24163450 0.07279700 3.319292 9.024598e-04
```

```
## omega 0.33815662 0.06144724 5.503203 3.729535e-08
## alpha1 0.83482811 0.24295106 3.436199 5.899382e-04
```

- (a) The ARMA/GARCH model is fitted for Del.Tbill. The model that corresponds to R parameter names is:

$$X_t - \mu = \text{ar1} \cdot (X_{t-1} - \mu) + a_t$$

$$a_t = \sigma_t \epsilon_t$$

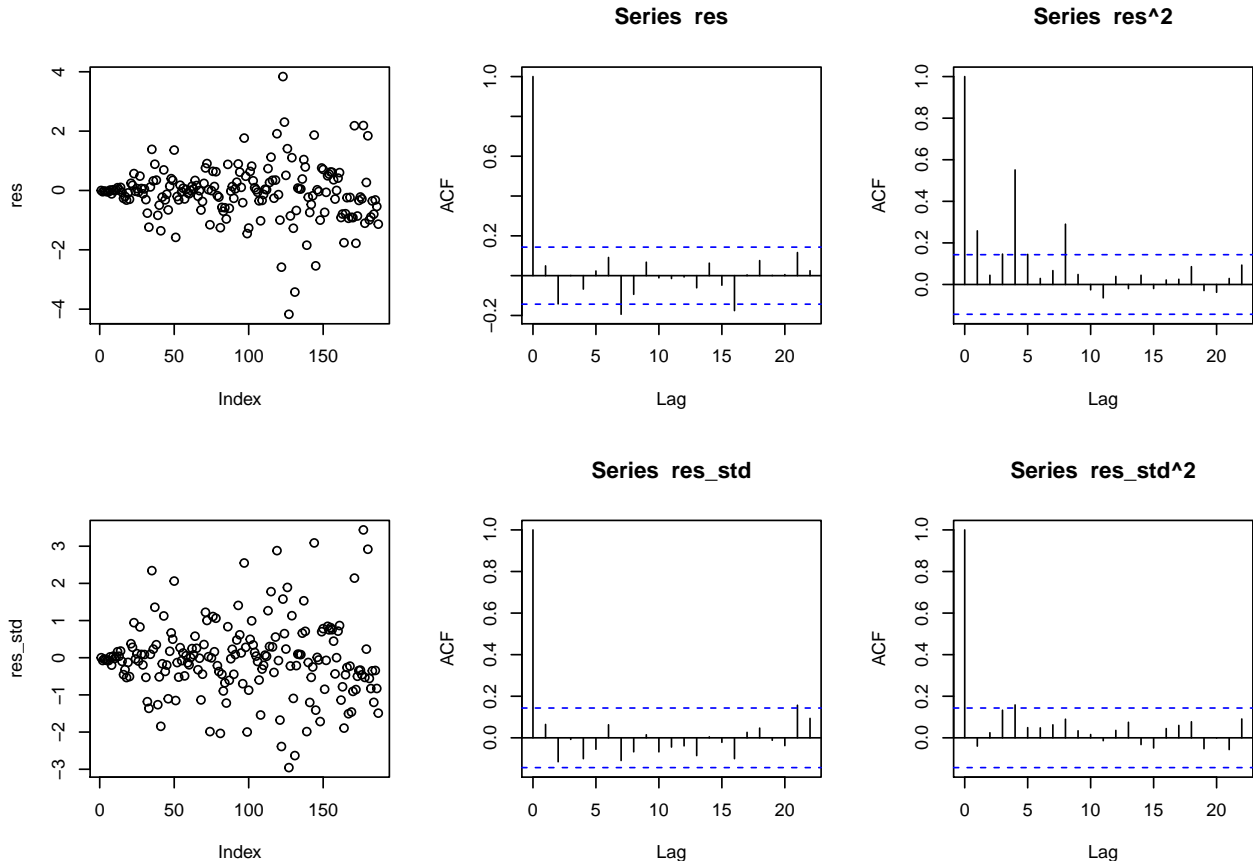
$$\sigma_t^2 = \text{omega} + \text{alpha1} \cdot a_{t-1}^2$$

$$\epsilon_t \sim \text{i.i.d white noise, mean 0, unit variance}$$

- (b) The parameter estimates are shown in the `summary(garch.model.Tbill)`, where we have: `mu = 0.08350`, `ar1 = 0.24163`, `omega = 0.33816`, and `alpha1 = 0.83483`.

Problem 3

```
res = residuals(garch.model.Tbill)
res_std = res / garch.model.Tbill@sigma.t
par(mfrow=c(2,3))
plot(res)
acf(res)
acf(res^2)
plot(res_std)
acf(res_std)
acf(res_std^2)
```



Answers (I'd like to answer (d) first, since (c) is actually based upon (d))

- (d): `garch.model.Tbill@sigma.t` contains the estimates of the conditional standard deviation of a_t from the GARCH model. Therefore, `res_std` is the realizations of $\epsilon_t = a_t/\sigma_t$.
- (a): `acf(res)` shows there is no significant autocorrelation in the residuals, which implies that the model fit of the conditional mean (AR(1)) is adequate.
- (b): `acf(res^2)` shows significant autocorrelation in the squared residuals a_t^2 , which implies a conditional heteroscedasticity - the reason why we are fitting the GARCH part.
- (c): `acf(res_std^2)` shows no significant autocorrelation in the standardized squared residuals $\epsilon_t^2 = (a_t/\sigma_t)^2$, which is consistent to the white noise assumption about ϵ_t , and implies that GARCH(1,0) fit for the conditional standard deviation is again adequate.
- (e): The time series plot of `res_std` shows the empirical distribution of `res_std` have heavier tail than normal distribution - there are quite a few points with absolute value > 1.96 , i.e. beyond 97.5% two-sided percentile of standard normal distribution.

Problem 4

```
# fit the model
garch.model.log.Tbill = garchFit(formula= ~arma(1,0) + garch(1,0), diff(log(Tbill)))

##
## Series Initialization:
## ARMA Model:          arma
## Formula Mean:        ~ arma(1, 0)
## GARCH Model:         garch
## Formula Variance:    ~ garch(1, 0)
## ARMA Order:          1 0
## Max ARMA Order:      1
## GARCH Order:         1 0
## Max GARCH Order:     1
## Maximum Order:       1
## Conditional Dist:     norm
## h.start:             2
## llh.start:           1
## Length of Series:     187
## Recursion Init:       mci
## Series Scale:         0.1506804
##
## Parameter Initialization:
## Initial Parameters:   $params
## Limits of Transformations: $U, $V
## Which Parameters are Fixed? $includes
## Parameter Matrix:
##           U           V      params includes
## mu      -0.63215833   0.6321583 0.05810348    TRUE
## ar1      -0.99999999   1.0000000 0.29238716    TRUE
## omega    0.00000100 100.0000000 0.10000000    TRUE
## alpha1   0.00000001   1.0000000 0.10000000    TRUE
## gamma1  -0.99999999   1.0000000 0.10000000    FALSE
## delta    0.00000000   2.0000000 2.00000000    FALSE
## skew     0.10000000 10.0000000 1.00000000    FALSE
## shape    1.00000000 10.0000000 4.00000000    FALSE
## Index List of Parameters to be Optimized:
```

```

##      mu      ar1  omega alpha1
##      1      2      3      4
## Persistence:                0.1
##
##
## --- START OF TRACE ---
## Selected Algorithm: nlminb
##
## R coded nlminb Solver:
##
## 0:      540.68757: 0.0581035 0.292387 0.100000 0.100000
## 1:      259.25047: 0.0581586 0.351171 1.04735 0.414755
## 2:      258.16987: 0.0582028 0.456570 1.03558 0.402981
## 3:      252.49959: 0.0574411 0.545312 0.727946 0.684987
## 4:      250.06907: 0.0574976 0.428738 0.487576 0.613163
## 5:      249.49697: 0.0730634 0.539003 0.542145 0.585597
## 6:      249.37315: 0.0770830 0.479971 0.564196 0.647037
## 7:      249.26377: 0.0814361 0.486281 0.508052 0.660433
## 8:      249.22636: 0.0809935 0.488008 0.527309 0.622664
## 9:      249.22557: 0.0800861 0.491156 0.531795 0.627067
## 10:     249.22447: 0.0791458 0.490668 0.531751 0.621237
## 11:     249.22437: 0.0788238 0.489868 0.531039 0.621091
## 12:     249.22436: 0.0788350 0.490084 0.531231 0.620527
## 13:     249.22436: 0.0788424 0.490071 0.531222 0.620629
## 14:     249.22436: 0.0788416 0.490071 0.531221 0.620624
##
## Final Estimate of the Negative LLH:
## LLH: -104.6908      norm LLH: -0.5598439
##      mu      ar1      omega      alpha1
## 0.01187989 0.49007057 0.01206114 0.62062414
##
## R-optimhess Difference Approximated Hessian Matrix:
##      mu      ar1      omega      alpha1
## mu      -12197.74425 -168.269073   -7034.6225    77.646879
## ar1      -168.26907 -234.866231    -192.8231    -6.463916
## omega    -7034.62251 -192.823054  -371223.5397  -1802.624159
## alpha1    77.64688   -6.463916   -1802.6242   -32.483982
## attr("time")
## Time difference of 0.009238958 secs
##
## --- END OF TRACE ---
##
##
## Time to Estimate Parameters:
## Time difference of 0.03180599 secs
summary(garch.model.log.Tbill)

##
## Title:
## GARCH Modelling
##
## Call:
## garchFit(formula = ~arma(1, 0) + garch(1, 0), data = diff(log(Tbill)))
##

```



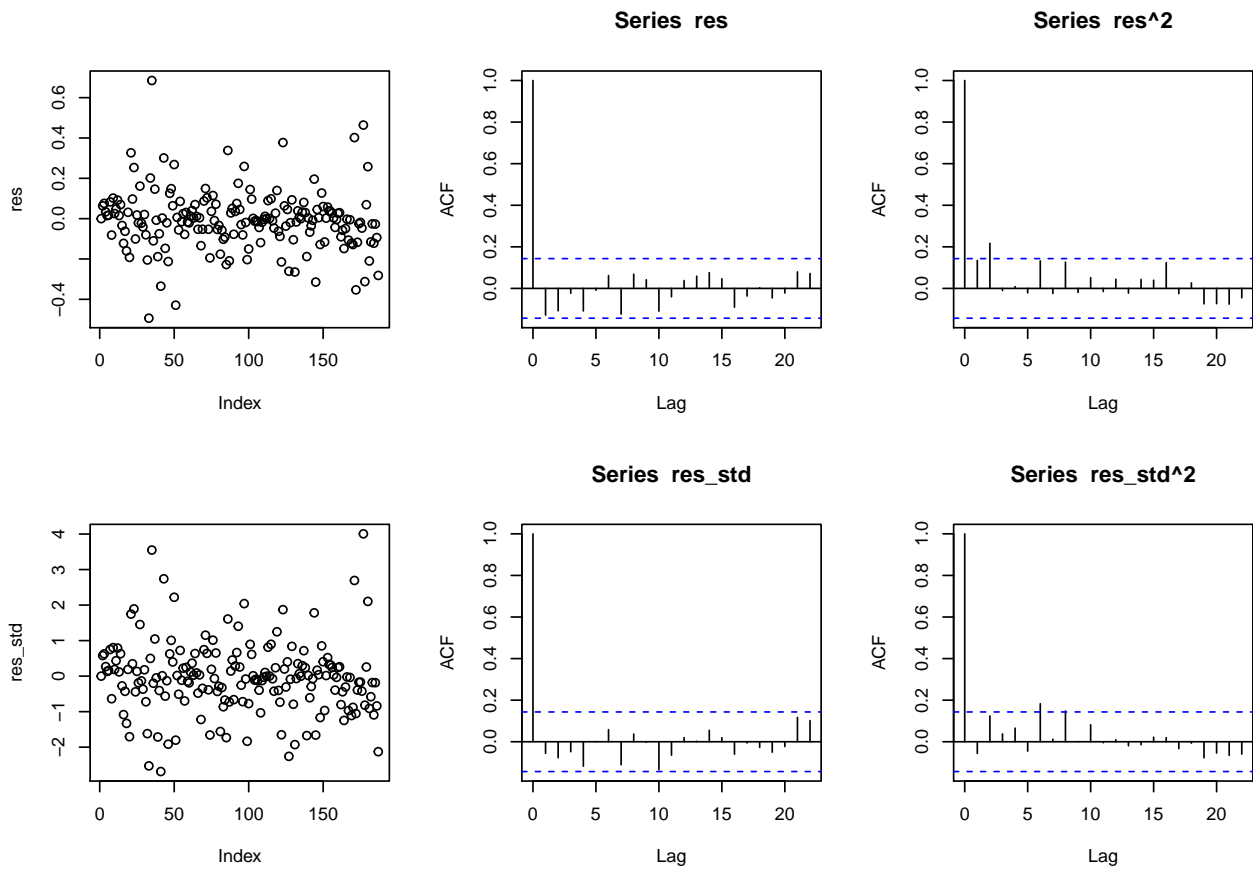
```

## Mean and Variance Equation:
## data ~ arma(1, 0) + garch(1, 0)
## <environment: 0x7fadcd0e108>
## [data = diff(log(Tbill))]
##
## Conditional Distribution:
## norm
##
## Coefficient(s):
##      mu      ar1      omega      alpha1
## 0.011880 0.490071 0.012061 0.620624
##
## Std. Errors:
## based on Hessian
##
## Error Analysis:
##      Estimate Std. Error t value Pr(>|t|)
## mu      0.011880   0.009373   1.267 0.20500
## ar1     0.490071   0.065891   7.438 1.03e-13 ***
## omega   0.012061   0.001961   6.150 7.76e-10 ***
## alpha1  0.620624   0.210910   2.943 0.00325 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Log Likelihood:
## 104.6908      normalized: 0.5598439
##
## Description:
## Mon Oct  1 15:38:54 2018 by user:
##
##
## Standardised Residuals Tests:
##
##                               Statistic p-Value
## Jarque-Bera Test      R      Chi^2 45.79453 1.137217e-10
## Shapiro-Wilk Test     R      W      0.9540278 9.29422e-06
## Ljung-Box Test        R      Q(10) 11.76569 0.3010438
## Ljung-Box Test        R      Q(15) 13.38437 0.5726355
## Ljung-Box Test        R      Q(20) 14.94135 0.7797538
## Ljung-Box Test        R^2 Q(10) 17.25489 0.06891087
## Ljung-Box Test        R^2 Q(15) 17.49449 0.2901728
## Ljung-Box Test        R^2 Q(20) 19.69902 0.4768933
## LM Arch Test          R      TR^2 12.22925 0.4274482
##
## Information Criterion Statistics:
##      AIC      BIC      SIC      HQIC
## -1.076907 -1.007792 -1.077797 -1.048902


# diagnostic plots
res = residuals(garch.model.log.Tbill)
res_std = res / garch.model.log.Tbill@sigma.t
par(mfrow=c(2,3))
plot(res)
acf(res)
acf(res^2)

```

```
plot(res_std)
acf(res_std)
acf(res_std^2)
```



```
par(mfrow=c(1,2))
plot(garch.model.Tbill@sigma.t)
lines(garch.model.Tbill@sigma.t)
plot(garch.model.log.Tbill@sigma.t)
lines(garch.model.log.Tbill@sigma.t)
```



Comments

- The main difference between the two is the conditional standard deviation series. Shown in the time series plot above.
- The σ_t of the log difference model is more “stable”, and the spike associated with large Tbill value around early 1980s also disappeared in the log transformed model.

Ford Data


(a)

```
library(quantmod)

## Loading required package: xts
## Loading required package: zoo
##
## Attaching package: 'zoo'
## The following object is masked from 'package:timeSeries':
##
##     time<-
## The following objects are masked from 'package:base':
##
##     as.Date, as.Date.numeric
## Loading required package: TTR
##
## Attaching package: 'TTR'
## The following object is masked from 'package:fBasics':
##
##     volatility
## Version 0.4-0 included new data defaults. See ?getSymbols.
FORD = getSymbols('F', from='1995-1-1', to='2017-12-31', auto.assign = F)

## 'getSymbols' currently uses auto.assign=TRUE by default, but will
## use auto.assign=FALSE in 0.5-0. You will still be able to use
## 'loadSymbols' to automatically load data. getOption("getSymbols.env")
## and getOption("getSymbols.auto.assign") will still be checked for
## alternate defaults.
##
## This message is shown once per session and may be disabled by setting
## options("getSymbols.warning4.0"=FALSE). See ?getSymbols for details.
##
## WARNING: There have been significant changes to Yahoo Finance data.
## Please see the Warning section of '?getSymbols.yahoo' for details.
##
## This message is shown once per session and may be disabled by setting
## options("getSymbols.yahoo.warning"=FALSE).
```

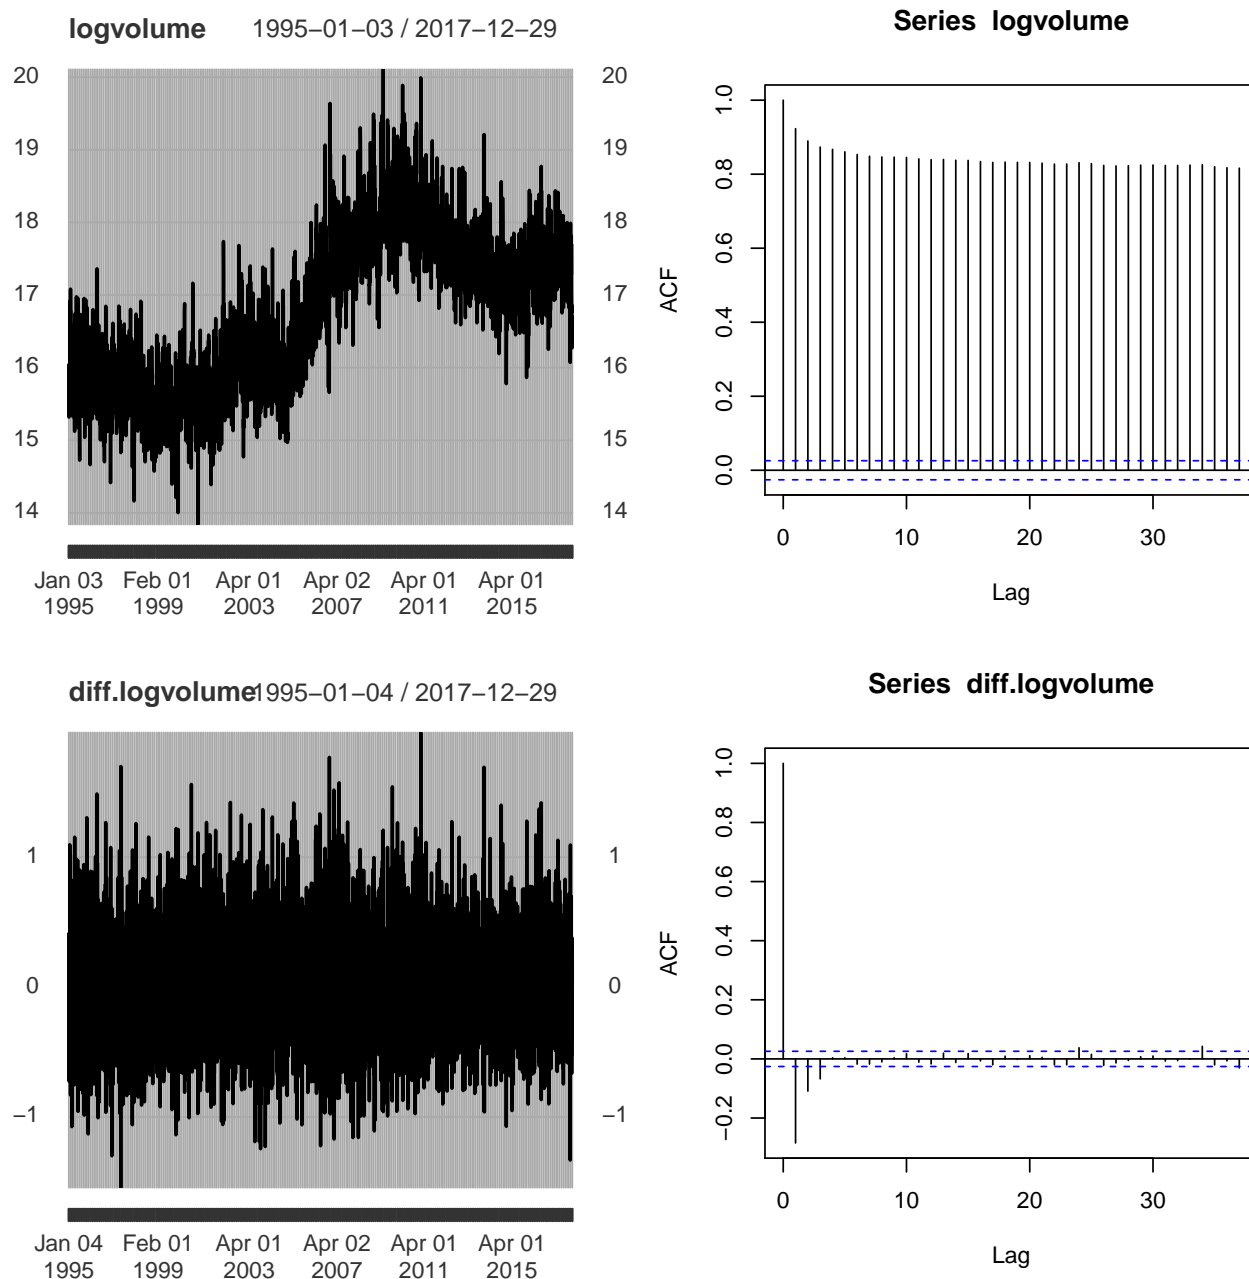
```
volume = FORD$F.Volume
plot(volume)
```



The volume has an extremely unstable size of variability. In particular, the variability of the series spikes among 2009 - 2011. Therefore, it would be a good idea to apply Box-Cox transformation to stabilize its variability.

(b)

```
logvolume = log(volume)
diff.logvolume = diff(logvolume)
diff.logvolume = diff.logvolume[!is.na(diff.logvolume)]
par(mfrow=c(2,2))
plot(logvolume)
acf(logvolume)
plot(diff.logvolume)
acf(diff.logvolume)
```



- The `logvolume` series has a clear trend in it, and the ACF decays slowly. It's not stationary.
- The first difference of `logvolume` exhibits a more stable behavior, and the ACF of it decays very fast. So we believe `diff.logvolume` is stationary.

(c)


```
arima.fit = auto.arima(diff.logvolume, ic='aicc', approximation=F, step=F)
summary(arima.fit)
```

```
## Series: diff.logvolume
## ARIMA(1,0,4) with zero mean
##
```

```
## Coefficients:
##          ar1      ma1      ma2      ma3      ma4
##          0.8598 -1.3442  0.2199  0.0650  0.067
## s.e.    0.0248   0.0284  0.0245  0.0221  0.017
##
## sigma^2 estimated as 0.1242:  log likelihood=-2176.21
## AIC=4364.42   AICc=4364.43   BIC=4404.4
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.004279167 0.352317 0.2698948 39.7425 277.3254 0.5384615
##              ACF1
## Training set -0.0001589462
```

Auto arima with AICC criteria finds optimal model ARIMA(1, 0, 4). The parameter estimates are shown in the summary above.

```
res = arima.fit$residuals
par(mfrow=c(1,2))
acf(res)
acf(res^2)
```



- `acf(res)` shows there is no significant autocorrelation in the residuals, which implies that the model fit of the conditional mean ARIMA(1, 0, 4) is adequate.
- `acf(res^2)` shows some autocorrelation in the squared residuals a_t^2 at lag=1, which implies a conditional heteroscedasticity. We can try GARCH fit to this part.

```
garch.model.logvolume = garchFit(
  formula= ~arma(1,4) + garch(1,1), diff.logvolume)
```

```
garch.model.logvolume@fit$coef
```

```
##          mu          ar1          ma1          ma2          ma3
## 6.853411e-05 5.257637e-01 -1.000000e+00 5.024231e-02 2.999536e-03
##          ma4          omega          alpha1          beta1
## -7.899813e-03 1.931606e-04 3.749038e-03 9.946603e-01
```

The fitted parameters of ARIMA(1,0,4)/GARCH(1,1) are shown in the list above.

```
arima.fit$aic
```

```
## [1] 4364.42
```

```
(garch.model.logvolume@fit$ics) * length(diff.logvolume)
```

```
##      AIC      BIC      SIC      HQIC
```


```
## 4363.129 4423.104 4363.101 4383.994
```

The ARIMA(1,0,4) conditional mean model alone has an AIC of 4364.42, while the ARIMA(1,0,4)/GARCH(1,1) model has a smaller AIC of 4363.129. Thus AIC implies that including the GARCH(1,1) component is useful indeed.

```
qqnorm(garch.model.logvolume@residuals)
```

```
qqline(garch.model.logvolume@residuals)
```

Normal Q-Q Plot



The QQ plot suggests that the residuals from the ARIMA/GARCH fit is still heavier than normal.

```
aparch.model.logvolume = garchFit(  
  formula= ~arma(1,4) + aparch(1,1), diff.logvolume)
```

```
## Warning in sqrt(diag(fit$cvar)): NaNs produced
```

```
aparch.model.logvolume@fit$coef
```

```
##      mu      ar1      ma1      ma2      ma3  
## 0.0000529336 0.5255967948 -0.9999999900 0.0503433988 0.0028967156  
##      ma4      omega      alpha1      gamma1      beta1  
## -0.0075848452 0.0003364403 0.0053547278 0.1443256624 0.9930301131  
##      delta  
## 1.7692790767
```

```
arima.fit$aic
```

```
## [1] 4364.42
```

```
(garch.model.logvolume@fit$ics) * length(diff.logvolume)
```

```
##      AIC      BIC      SIC      HQIC
```

```
## 4363.129 4423.104 4363.101 4383.994
```

```
(aparch.model.logvolume@fit$ics) * length(diff.logvolume)
```

```
##      AIC      BIC      SIC      HQIC
```

```
## 4371.712 4445.014 4371.670 4397.213
```

The APARCH(1,1) model has a larger AIC of 4371.712. Thus the AIC does not justify this additional complexity. The ARIMA/GARCH model is still the preferable one.